

Introduction to R – Transforming data, creating graphics

Jana Jarecki, Stephan Lewandowsky

11 July 2022

Agenda: What today teaches or repeats

Data Reading

- How to read data
- How to change column classes

Data Management

- How to manipulate strings

- How to transform data
- When summarize data by groups

Data Visualization

- The principles of graphics (*ggplot*)
- Two useful plots for modelers

A large package collection

The tidyverse collection

- A package collection, installing which downloads currently 25 packages at once (like MS Office that comes with Word, Excel, Powerpoint, etc.)
- `library(tidyverse)`, as of today, calls `library` for eight of those packages:
 - [readr](#) to read and write data
 - [tibble](#) to represent data.frames easier
 - [dplyr](#) to manipulate data e.g. long-to-wide and vice versa
 - [ggplot2](#) to plot graphic
 - [purrr](#) to apply functions
 - [tidyr](#) to tidy data
 - [stringr](#) to work with strings
 - [forcats](#) to deal with factors

R Tip: Reproducible Package Management

A usability problem

- `library(aaa)` can load package *aaa* only if it is installed
- Poor usability for you who want to execute my script
- I want a command that
 - checks: are package installed? not, installs them
 - and: Loads packages

A solution *pacman*

- `pacman::p_load(aaa, bbb)` checks: are packages *aaa, bbb*, installed? If not, installs them
- And: Loads the packages automatically

```
# delete this line afterwards!  
install.packages("pacman")  
# pacman:: calls a function in pacman  
pacman::p_load(tidyverse)
```

Note the notation `pacman::`

Reading data

Reading data with `readr()`

Setup

```
pacman::p_load(readr)  
setwd(dirname(rstudioapi::getSourceEditorContext()$path))
```

Data Reading

Read the data *ggplot-data.csv* with base R's command `read.csv()` and the tidyr's command `read_csv()`, what are the differences?

Reading Data

```
setwd(dir=dirname(rstudioapi::getSourceEditorContext())$path))
a <- read.csv("data/ggplot-data.csv")
a
```

##	id	condition	judgment	certainty	choice	model_X_pred	model_Y_pred
## 1	1	A	5.11287534	2.6285737	0	0.4540471	0.5397072
## 2	2	A	0.30590549	1.1001238	1	0.7245752	0.3362662
## 3	3	A	2.08938523	2.5760403	1	0.5282888	0.1851453
## 4	4	A	3.89858781	1.4530127	1	0.5845568	0.2678840
## 5	5	A	2.21280497	0.4742799	1	0.7198482	0.4388933
## 6	6	A	1.68162645	2.1691433	1	0.6927926	0.6610837
## 7	7	A	5.53456599	1.6955200	1	0.5637879	0.3150401
## 8	8	A	1.71602288	3.1873058	1	0.5598316	0.6052043
## 9	9	A	7.05527114	2.3796080	1	0.6997092	0.5334455
## 10	10	A	1.81185770	2.4180194	1	0.8762260	0.5504416
## 11	11	A	4.91460896	2.7048471	1	0.6336925	0.6543931
## 12	12	A	8.85993618	2.3881483	1	0.5753737	0.1014269
## 13	13	A	-3.16658210	2.3424111	1	0.6078865	0.1962579
## 14	14	A	1.16363370	2.2756260	1	0.7122976	0.3395638
## 15	15	A	0.60003599	1.6097678	1	0.5085801	0.5051917
## 16	16	A	3.90785119	2.4581210	1	0.8195041	0.3882232
## 17	17	A	0.14724124	2.1087371	0	0.3298549	0.6271167
## 18	18	A	-5.96936626	0.2959596	1	0.7159193	0.2901701
## 19	19	A	-6.32140079	-2.8573702	1	0.7284936	0.5888546

Reading Data

```
pacman::p_load(readr)
setwd(dirname(rstudioapi::getSourceEditorContext()$path))
d <- read_csv("data/ggplot-data.csv")
d
```

```
## # A tibble: 40 x 11
```

```
##       id condition judgment certainty choice model_X_pred model_Y_pred
##   <dbl> <chr>      <dbl>      <dbl> <dbl>      <dbl>      <dbl>
## 1     1     1 A        5.11      2.63     0        0.454      0.540
## 2     2     2 A        0.306     1.10     1        0.725      0.336
## 3     3     3 A        2.09      2.58     1        0.528      0.185
## 4     4     4 A        3.90      1.45     1        0.585      0.268
## 5     5     5 A        2.21      0.474    1        0.720      0.439
## 6     6     6 A        1.68      2.17     1        0.693      0.661
## 7     7     7 A        5.53      1.70     1        0.564      0.315
## 8     8     8 A        1.72      3.19     1        0.560      0.605
## 9     9     9 A        7.06      2.38     1        0.700      0.533
## 10    10    10 A        1.81      2.42     1        0.876      0.550
## # ... with 30 more rows, and 4 more variables: model_X_evidence <dbl>,
## #   model_Y_evidence <dbl>, alpha <dbl>, delta <dbl>
```

Reading Data: Column Types

subject "id" are in 99% of all cognitive data not meaningful numbers but IDs

```
d <- read_csv("data/ggplot-data.csv",
              col_type = cols(
                id = col_character()))
```

d

```
## # A tibble: 40 x 11
```

```
##   id    condition judgment certainty choice model_X_pred model_Y_pred
##   <chr> <chr>         <dbl>    <dbl> <dbl>         <dbl>         <dbl>
##  1 1      A           5.11      2.63    0           0.454         0.540
##  2 2      A           0.306     1.10    1           0.725         0.336
##  3 3      A           2.09      2.58    1           0.528         0.185
##  4 4      A           3.90      1.45    1           0.585         0.268
##  5 5      A           2.21      0.474    1           0.720         0.439
##  6 6      A           1.68      2.17    1           0.693         0.661
##  7 7      A           5.53      1.70    1           0.564         0.315
##  8 8      A           1.72      3.19    1           0.560         0.605
##  9 9      A           7.06      2.38    1           0.700         0.533
## 10 10     A           1.81      2.42    1           0.876         0.550
## # ... with 30 more rows, and 4 more variables: model_X_evidence <dbl>,
## #   model_Y_evidence <dbl>, alpha <dbl>, delta <dbl>
```

String Manipulation

String manipulation

```
s <- c("ab_1", "xy_1", "uv_1")  
print(s)
```

```
## [1] "ab_1" "xy_1" "uv_1"
```

```
pacman::p_load(stringr)  
str_flatten(s)
```

```
## [1] "ab_1xy_1uv_1"
```

```
str_split_fixed(s, pattern="_", n=2)
```

```
##      [,1] [,2]  
## [1,] "ab" "1"  
## [2,] "xy" "1"  
## [3,] "uv" "1"
```

```
d
```

```
## # A tibble: 40 x 11
```

```
##   id      condition judgment certainty choice model_X_pred model_Y_pred
```

Exercise: A Small Meaningless String Manipulation

- Make a new variable in the data `d` that holds the condition and choice in one string.
- (Optional) Substitute the pseudo-numeric `id` variable in `d` by a string of randomly selected 3 characters such as “nwt” for each subject.

```
pacman::p_load(stringr)
str_flatten(s)
str_split_fixed(s, pattern="_", n=2)
```

R Tip: De-numerize your IDs

Problem:

- If I don't transform the ID to a character, it can be summed up, averaged, etc which makes no sense
- Code can assume ID = 1.5 can exist

Solution

- In my preprocessing I de-numerize IDs

```
pacman::p_load(stringr)
# letters is an inbuilt
# R object
d$id <- replicate(
  nrow(d),
  str_flatten(sample(letters, 3)))
```

Note an inbuilt vector: letters

Selecting and Transforming

Data transformations

Selecting parts of data

```
pacman::p_load(dplyr)
filter(d, id=="9")
```

```
## # A tibble: 1 x 11
##   id    condition judgment certainty choice model_X_pred model_Y_pred
##   <chr> <chr>         <dbl>     <dbl> <dbl>         <dbl>         <dbl>
## 1 9      A           7.06      2.38    1           0.700         0.533
## # ... with 4 more variables: model_X_evidence <dbl>, model_Y_evidence <dbl>,
## #   alpha <dbl>, delta <dbl>
```

```
filter(d, condition=="A", judgment > 6)
```

```
## # A tibble: 2 x 11
##   id    condition judgment certainty choice model_X_pred model_Y_pred
##   <chr> <chr>         <dbl>     <dbl> <dbl>         <dbl>         <dbl>
## 1 9      A           7.06      2.38    1           0.700         0.533
## 2 12     A           8.86      2.39    1           0.575         0.101
## # ... with 4 more variables: model_X_evidence <dbl>, model_Y_evidence <dbl>,
## #   alpha <dbl>, delta <dbl>
```

```
select(d, model_X_pred)
```

Transformation

Transforming data: wide to long

```
select(d, id, model_X_pred, model_Y_pred)[1:3,]
```

```
## # A tibble: 3 x 3
##   id    model_X_pred model_Y_pred
##   <chr>      <dbl>      <dbl>
## 1 1          0.454          0.540
## 2 2          0.725          0.336
## 3 3          0.528          0.185
```

We need a *long* format of this data:

```
## # A tibble: 80 x 3
##   id    model      pred
##   <chr> <chr>      <dbl>
## 1 1    model_X_pred 0.454
## 2 1    model_Y_pred 0.540
## 3 2    model_X_pred 0.725
## 4 2    model_Y_pred 0.336
## 5 3    model_X_pred 0.528
## 6 3    model_Y_pred 0.185
## 7 4    model_X_pred 0.585
## 8 4    model_Y_pred 0.268
## 9 5    model_X_pred 0.720
```

Transforming data: wide to long

```
pacman::p_load(tidyr)
dl <- pivot_longer(
  data = d,
  cols=contains("pred"),
  names_to = "model",
  values_to = "pred")
select(dl, id, model, pred, choice)
```

```
## # A tibble: 80 x 4
##   id     model      pred choice
##   <chr> <chr>      <dbl>  <dbl>
## 1 1     model_X_pred 0.454    0
## 2 1     model_Y_pred 0.540    0
## 3 2     model_X_pred 0.725    1
## 4 2     model_Y_pred 0.336    1
## 5 3     model_X_pred 0.528    1
## 6 3     model_Y_pred 0.185    1
## 7 4     model_X_pred 0.585    1
## 8 4     model_Y_pred 0.268    1
## 9 5     model_X_pred 0.720    1
## 10 5     model_Y_pred 0.439    1
## # ... with 70 more rows
```

Transforming data: long to wide

```
dw <- pivot_wider(  
  data = dl,  
  names_from = model,  
  values_from = pred)  
select(dw, id, model_X_pred, model_Y_pred)
```

```
## # A tibble: 40 x 3  
##   id    model_X_pred model_Y_pred  
##   <chr>      <dbl>      <dbl>  
## 1 1      0.454      0.540  
## 2 2      0.725      0.336  
## 3 3      0.528      0.185  
## 4 4      0.585      0.268  
## 5 5      0.720      0.439  
## 6 6      0.693      0.661  
## 7 7      0.564      0.315  
## 8 8      0.560      0.605  
## 9 9      0.700      0.533  
## 10 10     0.876      0.550  
## # ... with 30 more rows
```

Transforming data: summary by groups

- M judgment by condition?
- SD of judgments by condition?
- SE of judgments by condition?

```
d[1:5]
```

```
## # A tibble: 40 x 5
##   id    condition judgment certainty choice
##   <chr> <chr>      <dbl>      <dbl> <dbl>
## 1 1      A          5.11       2.63    0
## 2 2      A          0.306      1.10    1
## 3 3      A          2.09       2.58    1
## 4 4      A          3.90       1.45    1
## 5 5      A          2.21       0.474   1
## 6 6      A          1.68       2.17    1
## 7 7      A          5.53       1.70    1
## 8 8      A          1.72       3.19    1
## 9 9      A          7.06       2.38    1
## 10 10     A          1.81       2.42    1
## # ... with 30 more rows
```

Exercise: summary by groups

- Compute the M and SE of the choice for each condition and store it as `d_obs_msd`
- Compute the M of the predictions for each condition and for each model and store it in `d_pred`

```
# Code reminder  
dg <- group_by(d, condition)  
d_m <- summarise(dg, across(  
  .cols="judgment",  
  .fns=list(M=mean)))
```

Two Dialects in the R Language

This code is the same ...

```
dg <- group_by(d, condition)
d_m <- summarise(dg, across(
  .cols="judgment",
  .fns=list(M=mean)))
d_m
```

```
## # A tibble: 2 x 2
##   condition judgment_M
##   <chr>           <dbl>
## 1 A             2.08
## 2 B             0.687
```

... as this code

```
d_m <- d %>%
  group_by(condition) %>%
  summarise(across(
    .cols="judgment",
    .fns=list(M=mean)))
d_m
```

```
## # A tibble: 2 x 2
##   condition judgment_M
##   <chr>           <dbl>
## 1 A             2.08
## 2 B             0.687
```

Note the notation %>%

Two Dialects in the R Language

```
a %>% f(b)  
# will execute  
f(a, b)
```

%>% passes what's left of it automatically as **first argument** to the function right of it, before any arguments that we give the function

In this case a becomes the first argument of f

- %>% is called a pipe
- comes from the *magrittr* package by Stefan Milton Bache

Plotting

Plotting with *ggplot2*

- *ggplot2* is a powerful and flexible package for making and saving graphics by Hadley [Wickham \(and others\)](#)
- it follows a quite straightforward logic
- it produces (almost) any type of static and animated 2D graphics

```
setwd(dirname(rstudioapi::getSourceEditorContext())$path))  
pacman::p_load(ggplot2) # install.package(pacman)
```

Logic of Graphics

In principle, a plot ...



... involves three questions


- *point, line, bar, boxplot, errorbar, etc.*
Which geometric objects?
- *x-axis = ?, y-axis = ?, colors = ?, etc.*
Which aesthetic dimensions map which variables?
- *title, color scales, axis texts, ...*
E.g., color-scale/theme adjustments that don't map variables

A Simple Graphic

Implemented in *ggplot*

```
d <- read.csv("data/ggplot-data.csv")

ggplot(data=d) +
  geom_point(mapping=aes(
    x=judgment,
    y=certainty))
```



... involves three questions

- *point, line, bar, boxplot, errorbar, etc.*
Which geometric objects?
- *x-axis = ?, y-axis = ?, colors = ?, etc.*
Which aesthetic dimensions map which variables?
- *title, color scales, axis texts, ...*
Layout adjustments that do not map any variables

Exercise: Some Extra Simple Plots

- Plot points with the condition on the x-axis and judgment on the y-axis
- Plot boxplots of judgment with the condition on the x-axis


```
head(d, n=2)
```

```
##   id condition  judgment certainty choice model_X_pred model_Y_pred
## 1  1          A 5.1128753  2.628574      0    0.4540471    0.5397072
## 2  2          A 0.3059055  1.100124      1    0.7245752    0.3362662
##   model_X_evidence model_Y_evidence    alpha    delta
## 1         0.13571107         0.8642889 0.4244226 4.582245
## 2         0.09960571         0.9003943 0.3057628 1.385302
```


- (Advanced, optional) Plot the mean judgment by condition as single points. Add errorbars representing the SD to the plot. Make the errorbar whisker smaller. Plot bars of **choice** on the x-axis, and check what exactly do you see? Try out `geom_dotplot(x = judgment)`.

Solution Some Extra Simple Plots

```
ggplot(data=d) +  
  geom_point(mapping=aes(  
    x=condition,  
    y=judgment))
```



```
ggplot(data=d) +  
  geom_boxplot(mapping=aes(  
    x=condition,  
    y=judgment))
```




Exercise: Some More Simple Plots

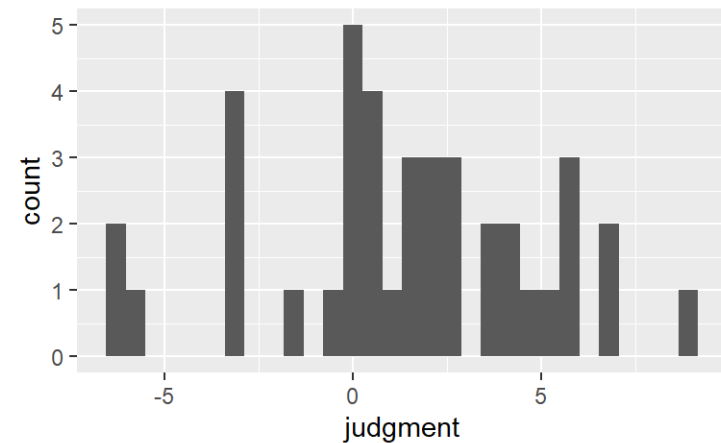
- Add `+geom_jitter(aes(x=condition, y=judgment))` with a `+` to the boxplot, see what it does
- Use a histogram as geometric object to show the distribution of judgment

Solution: Some Extra Simple Plots

```
ggplot(data=d) +  
  geom_boxplot(mapping=aes(  
    x=condition,  
    y=judgment)) +  
  geom_jitter(mapping=aes(  
    x=condition,  
    y=judgment))
```




```
ggplot(data=d) +  
  geom_histogram(mapping=aes(  
    x=judgment))
```



Changes the overall look


Get rid of this grey background

```
# Without a new theme being set
ggplot(data=d) +
  geom_point(mapping=aes(
    x=judgment,
    y=certainty))
```



A theme is set defining the look of *all subsequent* plots

```
theme_set(theme_classic()) # sets theme
ggplot(data=d) +
  geom_point(mapping=aes(
    x=judgment,
    y=certainty))
```



Note the (): `theme_classic()`


More themes:

<https://jrnold.github.io/ggthemes/>

Aesthetic mapping to dimensions beyond x and y

This plot ...

```
ggplot(data=d) +  
  geom_point(mapping=aes(  
    x=judgment,  
    y=certainty,  
    color=condition)  
  )
```



... shows more than two variables


- *x- and y-axes* are the core aesthetic dimensions mapping variables
`aes(x=judgment, y=certainty)`
- But they can map only two variables
- Beyond the axes, *color, shapes, linetypes, etc* are additional aesthetic dimensions mapping variables
`aes(color=condition)`

Exercise: Mapping > 2 variables

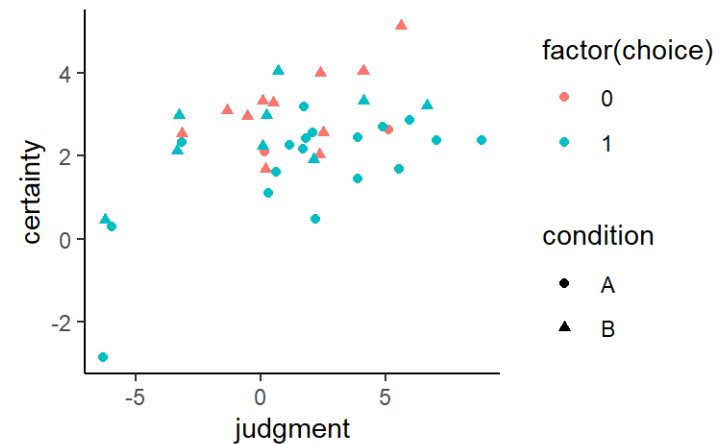
- Show the condition by point shape, rather than than color
- Show the variable choice by color
- (Optional) Change the shape of the points to a cross and a triangle, change the color scale to be grey to white, change the color scale to be discrete, change the title of the scales to “Observations” and “Experimental Condition”

Exercise: Mapping 2+ variables

```
ggplot(data=d) +  
  geom_point(mapping=aes(  
    x=judgment,  
    y=certainty,  
    shape=condition)  
  )
```



```
ggplot(data=d) +  
  geom_point(mapping=aes(  
    x=judgment,  
    y=certainty,  
    shape=condition,  
    color=factor(choice)  
  )
```



Advanced Exercises

Exercise: Qualitative Prediction Plots


```
# Transformation
se <- function(x) {sd(x) / sqrt(length(x))}
obs_msd <- d %>%
  group_by(condition) %>%
  summarise(across(choice,
                    .fns=list(M=mean, SD=sd, SE=se)))

pred_msd <- d %>%
  group_by(condition) %>%
  summarise(
    across(ends_with("pred"),mean)) %>%
  pivot_longer(
    cols=contains("pred"),
    names_to = "model",
    values_to = "M")
N <- length(unique(d$id))
```

```
p <- ggplot(data=obs_msd, aes(
  x=condition,
  y=choice_M)) +
  geom_col(width=.25,
           color="black",
           fill="lightgrey") +
  geom_errorbar(aes(
    ymin=choice_M,
    ymax=choice_M + choice_SE),
    width=.1) +
  geom_point(data=pred_msd, aes(
    x=condition,
    y=M,
    shape=model),
    size=3.5)
# --> Look at p
```

Exercise: Qualitative Prediction Plots (cont.)

```
# More cosmetics to the plot
p + scale_y_continuous(expand=c(0,NA)) +
  scale_shape_discrete(name="Predictions", labels=c("Model X", "Model Y")) +
  ylab("Predicted/Observed Choice (M, SE)") +
  xlab("Condition") +
  theme(legend.position=c(.9,.9),
        legend.background=element_rect(color="black")) +
  ggtitle("Qualitative Model Comparison")
```



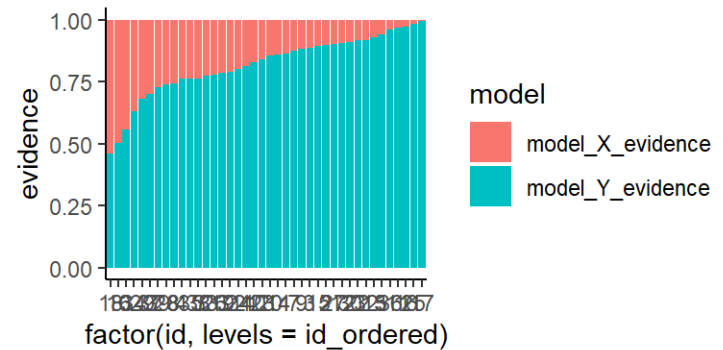
Advanced Exercise: Evidence Plots

```
# Transformation
dl <- pivot_longer(
  data = d,
  cols= c("model_X_evidence",
          "model_Y_evidence"),
  names_to = "model",
  values_to = "evidence")

id_order <- order(d$model_Y_evidence)
id_ordered <- d$id[id_order]
```

```
p <- ggplot(data=dl) +
  geom_col(mapping=aes(
    x=factor(id, levels=id_ordered),
    y=evidence,
    fill=model))

p
```



Advanced Exercise: Evidence Plots (cont.)

```
# more cosmetics
p + scale_y_continuous(expand=c(0,0)) +
  scale_fill_grey("Models", labels=c("Model X", "Model Y")) +
  theme(axis.text.x = element_blank(),
        axis.ticks.x = element_blank()) +
  ylab("Evidence Strength") +
  xlab("Participants") +
  ggtitle("Model Comparison")
```

